

# Transfer Learning

Utilizing Pre-trained Model for Better Results

Dr. Muhammad Sajjad

R.A: Kaleem Ullah

R.A: Imran Nawar


# Overview

- **Challenges in training custom Deep Learning models**
- **Transfer Learning**
- **How Transfer Learning Works**
- **Kinds of Transfer Learning**
- **Pre-trained Models**
- **ImageNet**
- **ImageNet Competition**
- **AlexNet Architecture**
- **LeNet5**
- **VGG16/19**
- **Problems with Very Deep Networks**
- **RESNET**
- **INCEPTON**
- **1x1 Convolution**
- **EFFICIENTNET**
- **MobileNet**
- **Applications of Transfer Learning**

# Challenges in Training Custom Deep Learning Models


## Data:

- Deep learning models typically require a large amount of labeled data to learn effectively.
- Gathering and labeling this data can be time-consuming and expensive.
- Without sufficient data, the model may not generalize well to new, unseen examples, leading to poor performance.



## Training Time:

- Training deep learning models can be computationally intensive and time-consuming.
- Depending on the complexity of the model architecture, size of the dataset, and available computational resources, training can take days, weeks, or even longer.
- Longer training times also increase the cost associated with experimentation and model development.



# Challenges in Training Custom Deep Learning Models

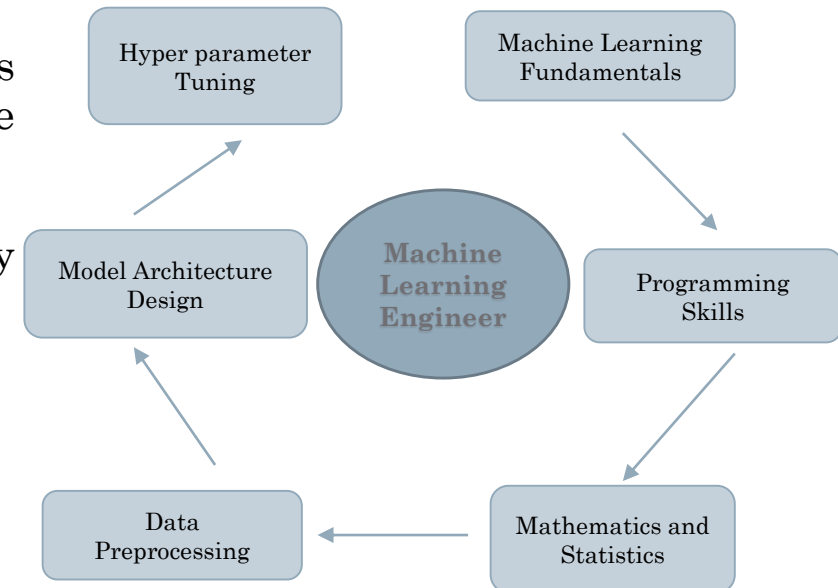
## Computational Resources:

- Training deep learning models often requires significant computational resources, including powerful GPUs or even specialized hardware like TPUs.
- Not everyone has access to these resources, limiting the ability to train complex models effectively.



## Expertise Requirement:

- Building and training custom deep learning models requires expertise in machine learning, deep learning, and software engineering.
- This expertise may not be readily available to everyone, especially those new to the field.



# Solution

# Transfer Learning

- A powerful technique in deep learning that allows us to reuse knowledge gained from solving one problem to tackle a different but related problem.
- Instead of training a model from scratch, we start with a pre-trained model and fine-tune it for the new task.
- It will not only speed up training considerably, but also requires significantly less training data.

## Need of Transfer Learning?

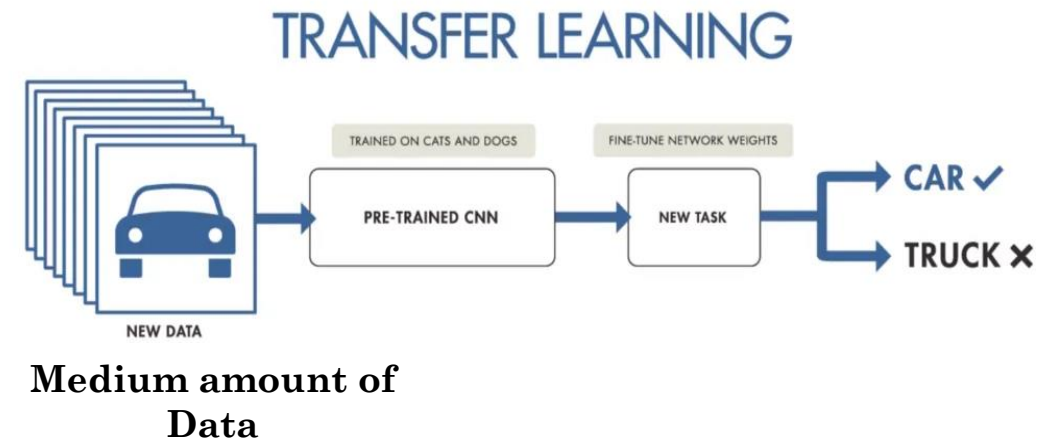
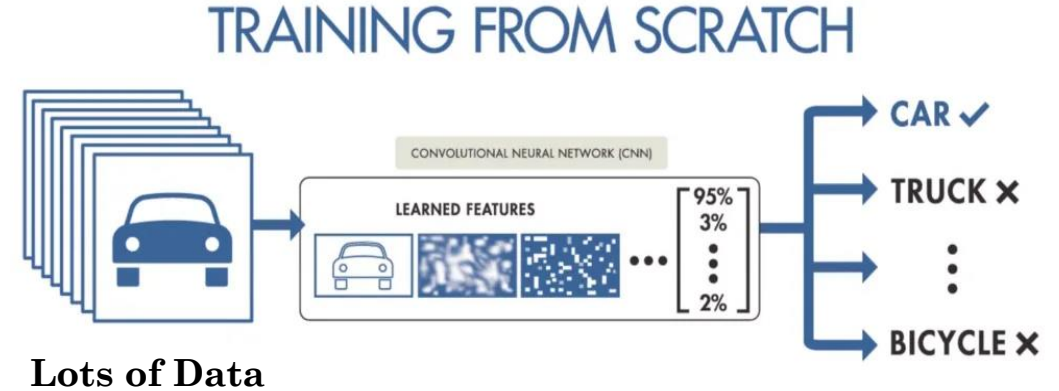
- If our dataset is really small
- Low Computation Power
- If our dataset is similar to pre-trained data then we have to only fine tuning our model it would save lot of time.

## Advantages:

- Work with Limited data
- Reduced training time
- Improved neural network performance(in most cases)

## Limitations

- Dataset is completely different from pre-trained data.




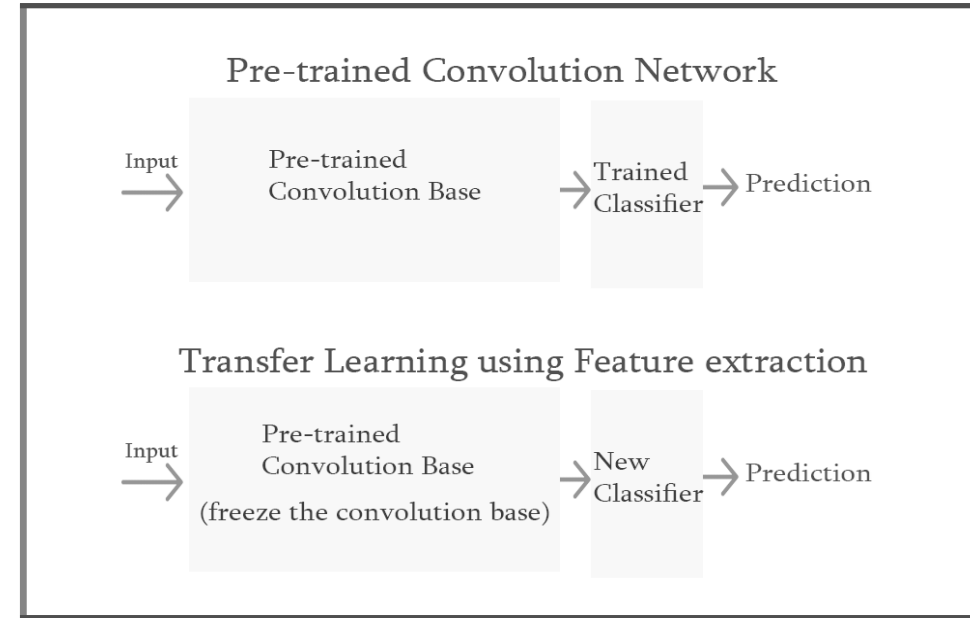
# How Transfer Learning Works

## Pre-trained Model:

- Begin with a model that has already been trained on a large dataset for a specific task (e.g., image classification using ImageNet).
- This pre-trained model has learned useful features and patterns from the data.

## Application to a New Task:

- Add new task-specific layers (e.g., an output layer) on top of the base model.
- Fine-tune the entire model on the new task using a smaller dataset.
- The base model's features are leveraged for the new task.



# Kinds of Transfer Learning

## **Feature Extraction:**

- Use the base model as a fixed feature extractor.
- Extract features from intermediate layers (e.g., convolutional layers in a CNN).
- Feed these features to a new classifier (e.g., a fully connected layer).
- Commonly used when data for the new task is limited.

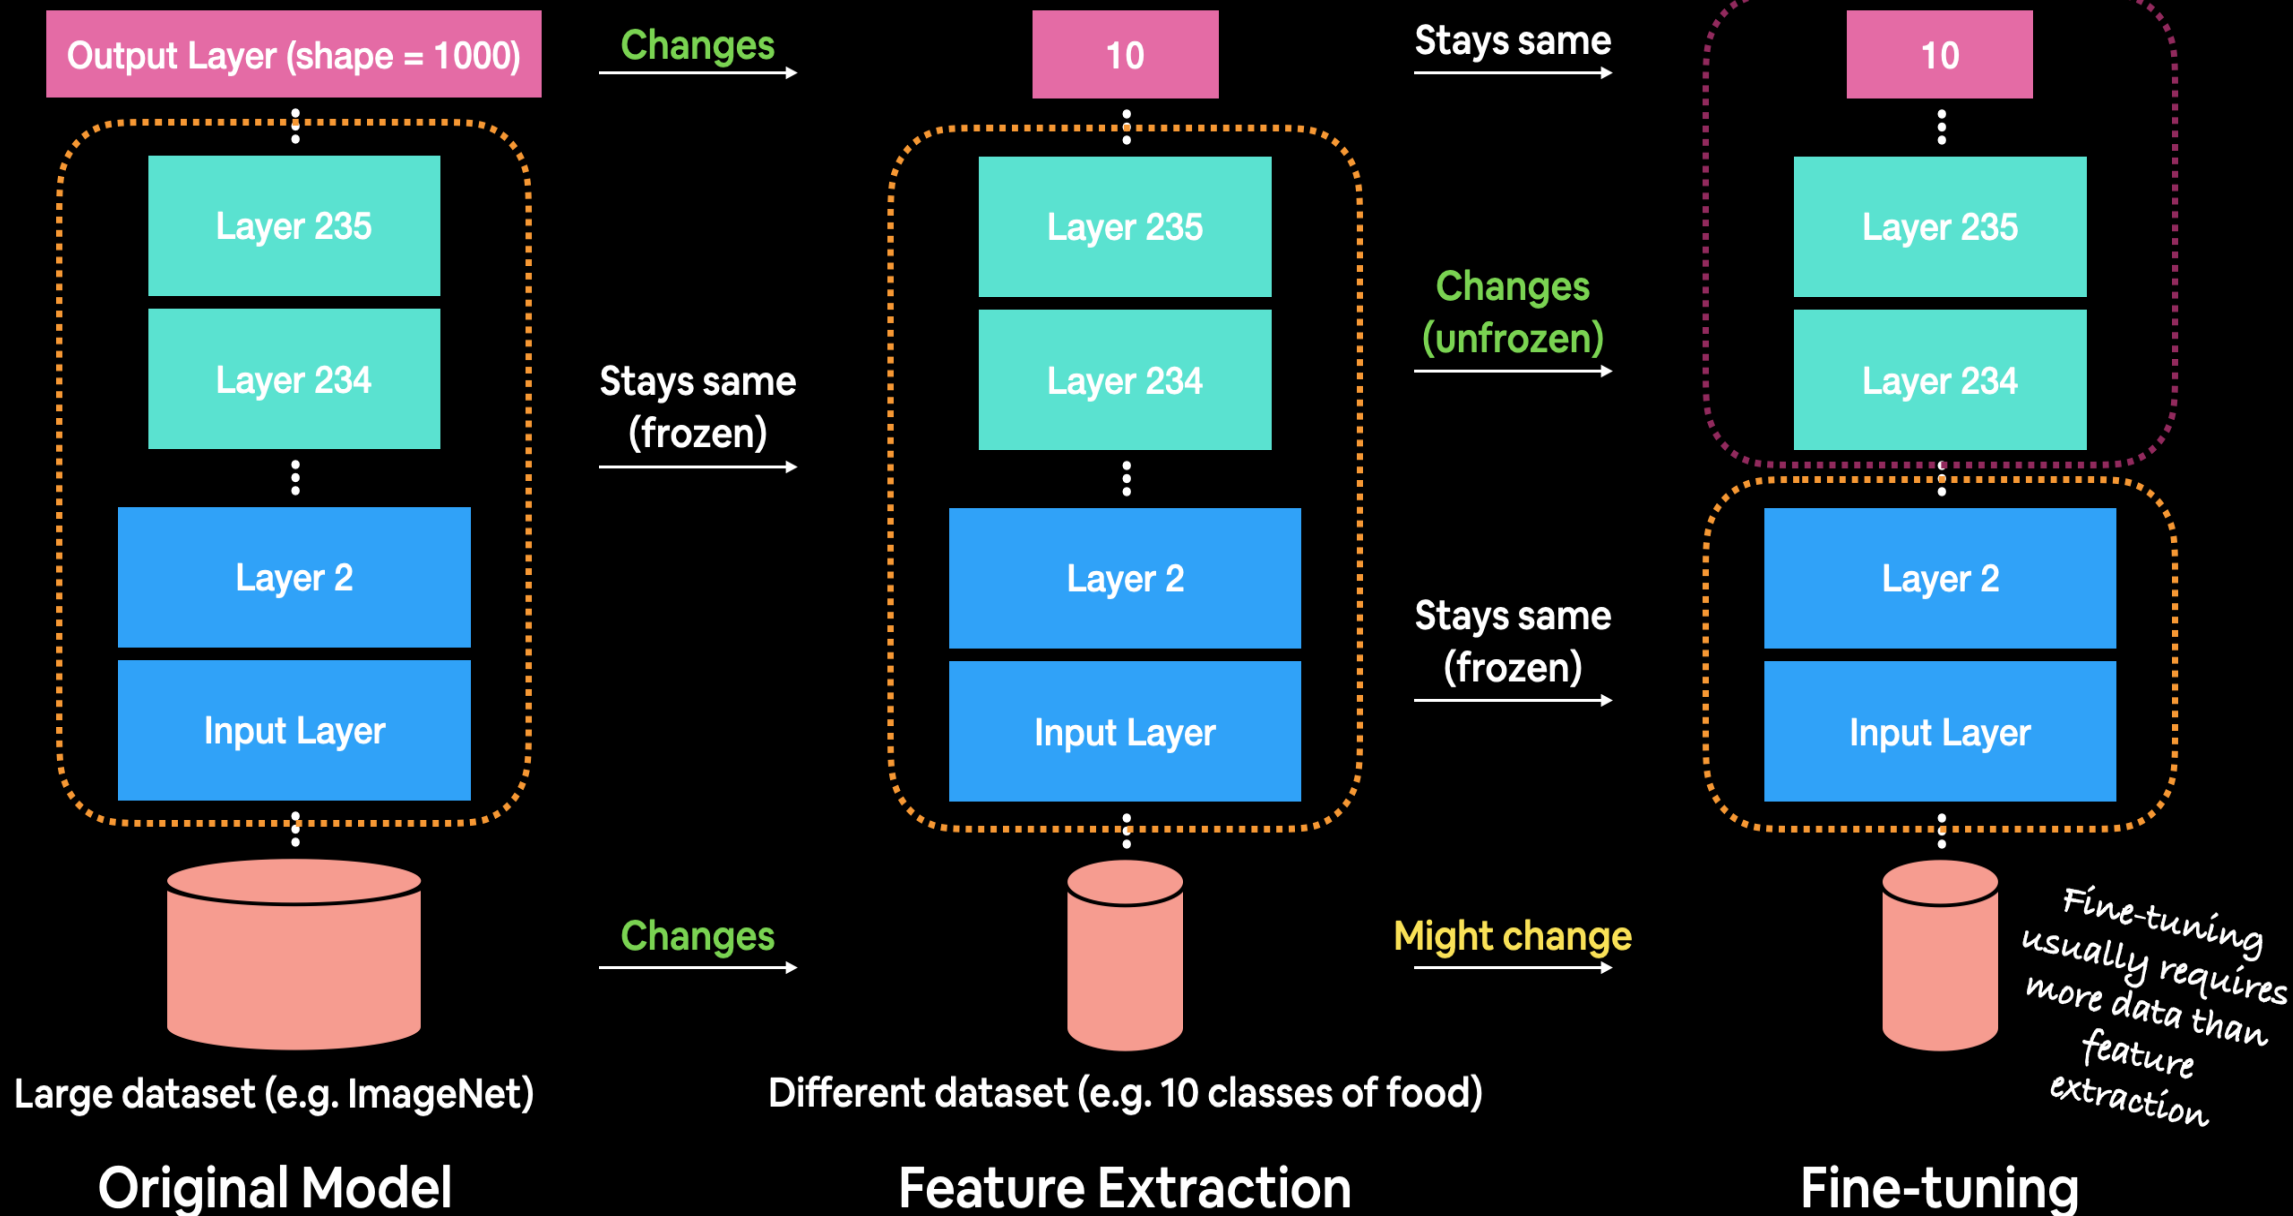
## **Fine-Tuning:**

- Fine-tune layer in the base model for the new task.
- Adjust the weights of the base model using the new task's data.
- Useful when the new task is closely related to the original task.



# Kinds of Transfer Learning

Top layers get trained on new data



← No gradient flow

○ Neurons from a pre-trained model

● Appended neurons

A diagram of a feedforward neural network with four layers. The first three layers consist of orange circular nodes, and the fourth layer consists of teal circular nodes. Dashed lines represent the connections between nodes in adjacent layers. Vertical ellipses (three dots) are used to indicate that there are multiple nodes in each layer, not just the ones explicitly drawn.

○ Full pre-trained network

- Appended network


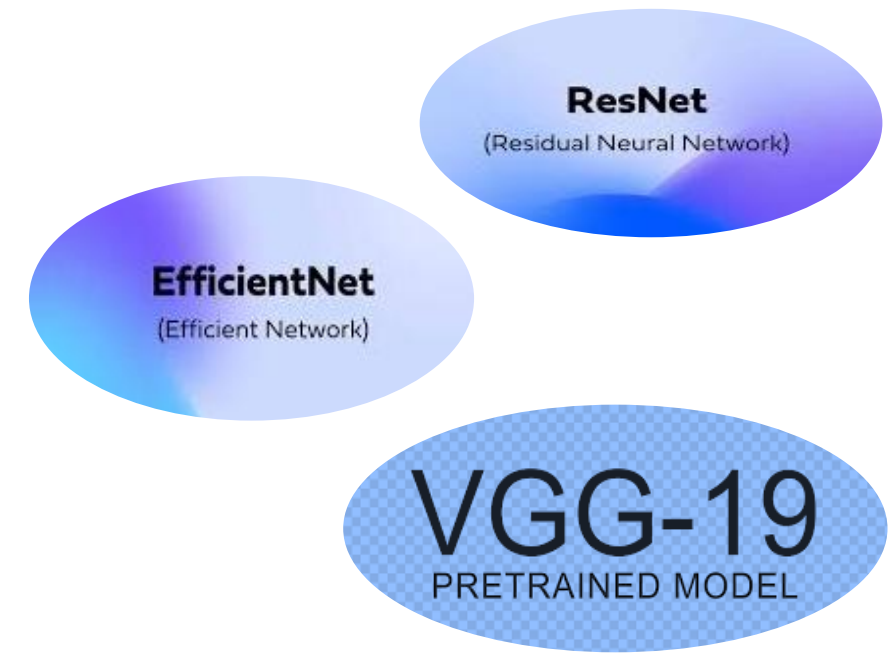
# Pre-trained Models

A **pre-trained model** is a deep learning model that has been trained on a large dataset and can be fine-tuned for a specific task.

- Pre-trained models serve as a starting point for developing deep learning models.
- They provide initial weights and biases that can be fine-tuned for specific tasks.
- Reuse lower layers of a pre-trained model for feature extraction, training only the final layers specific to your project.

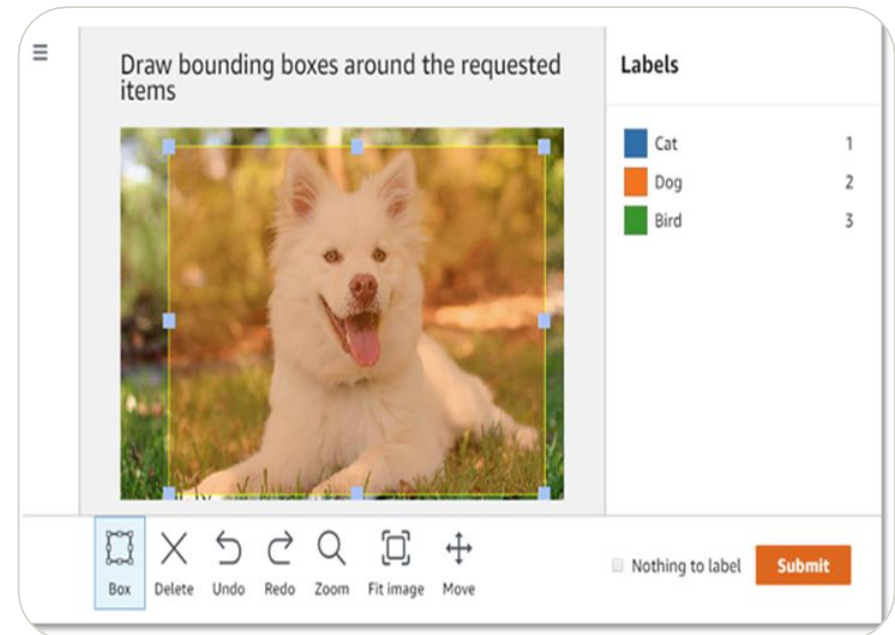
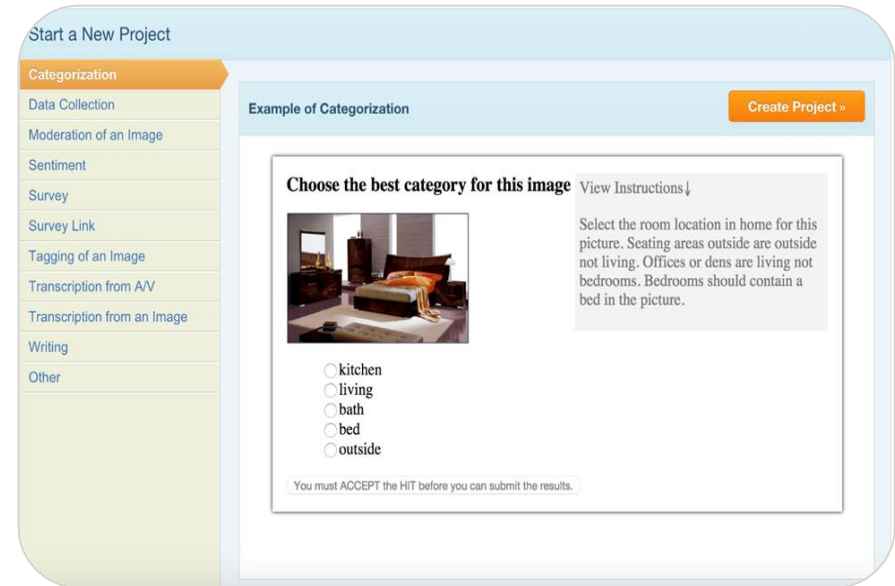
### Benefits:

- Save time and resources.
- Achieve higher accuracy with pre-learned features.
- Access models trained on large datasets (e.g., ImageNet with 14 million images).



# ImageNet

- The ImageNet dataset contained over **14 million** labeled images.
- The dataset contains more than 20,000 categories.
- A prominent computer scientist, **Fei-Fei Li**, co-founded the ImageNet project in 2009 (initiated in 2006) along with other researchers. Their work on ImageNet significantly advanced computer vision research and deep learning development.
- There are **1 million** Images with bounding box labelling as well for the purpose of **Object Localization** task, where the goal is to identify not only the object but also its precise location within the image.
- They used Amazon Mechanical Turk to help with the classification of images.
- Amazon Mechanical Turk (MTurk) is great for crowdsourcing tasks using images.





# ImageNet Competition


- The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) starts in 2010 , and the goal was to Highlight the best model for classification to the research Community.
- The dataset which were used in the challenge was a **subset** of ImageNet which consists of around **1.2 million** images from **1000 classes**.
- At first The peoples were use **ML Algorithms**.
- The Error rate was **28%** for the first time.
- In **2011** the error rate reduces to **25%** using ML algorithms.
- The Revolution in **2012** when **Geoffrey Hinton** Participated in this challenge with his CNN based Model **AlexNet**.
- It's effective implementation of deep learning algorithm, ReLU as an activation function and as well as uses GPU instead of CPU, significantly improved performance.
- AlexNet achieved a significantly lower error rate in the ILSVRC 2012 competition compared to previous approaches.



# Winning Models of the ImageNet Competition

YEAR	WINNER	TOP 5 ERROR RATE %
2012	ALEXNET	15.3
2013	ZFNET	11.2
2014	INCEPTION V1 (GoogLeNet) VGG NET (Runner up)	6.67 7.3
2015	ResNet	3.57
2016	ResNeXt	4.1
2017	SENet	2.251

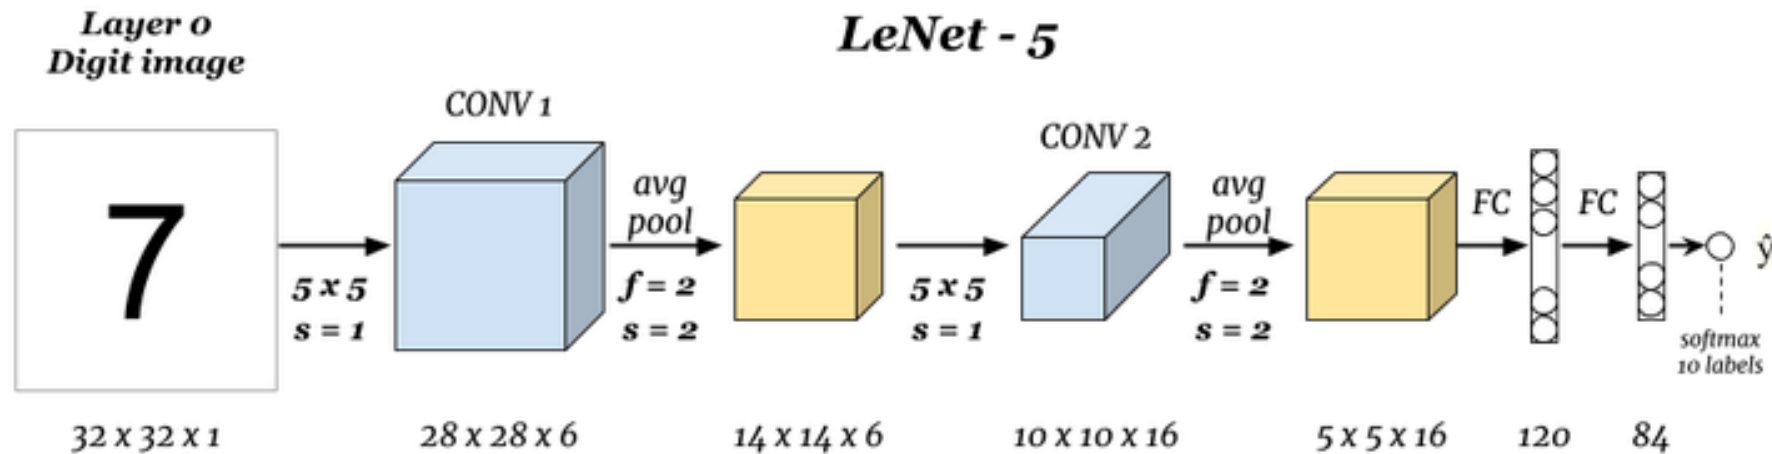
# AlexNet Architecture





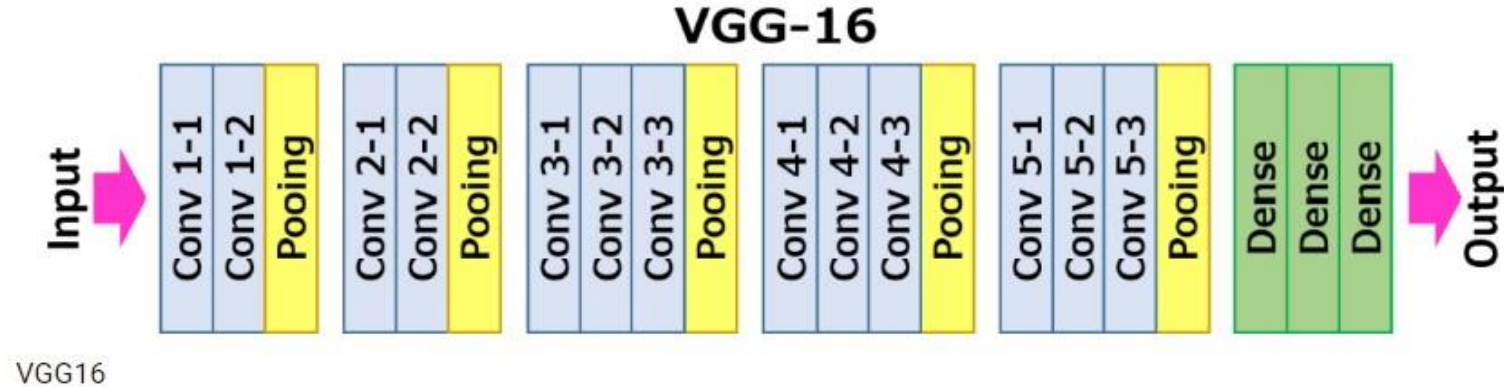
# LeNet5 Architecture

- LeNet-5 is the earliest CNN architecture, developed by Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner in 1998.
- It was primarily designed for handwritten digit recognition tasks, particularly recognizing digits in postal codes on letters



# VGG16/19 Architecture

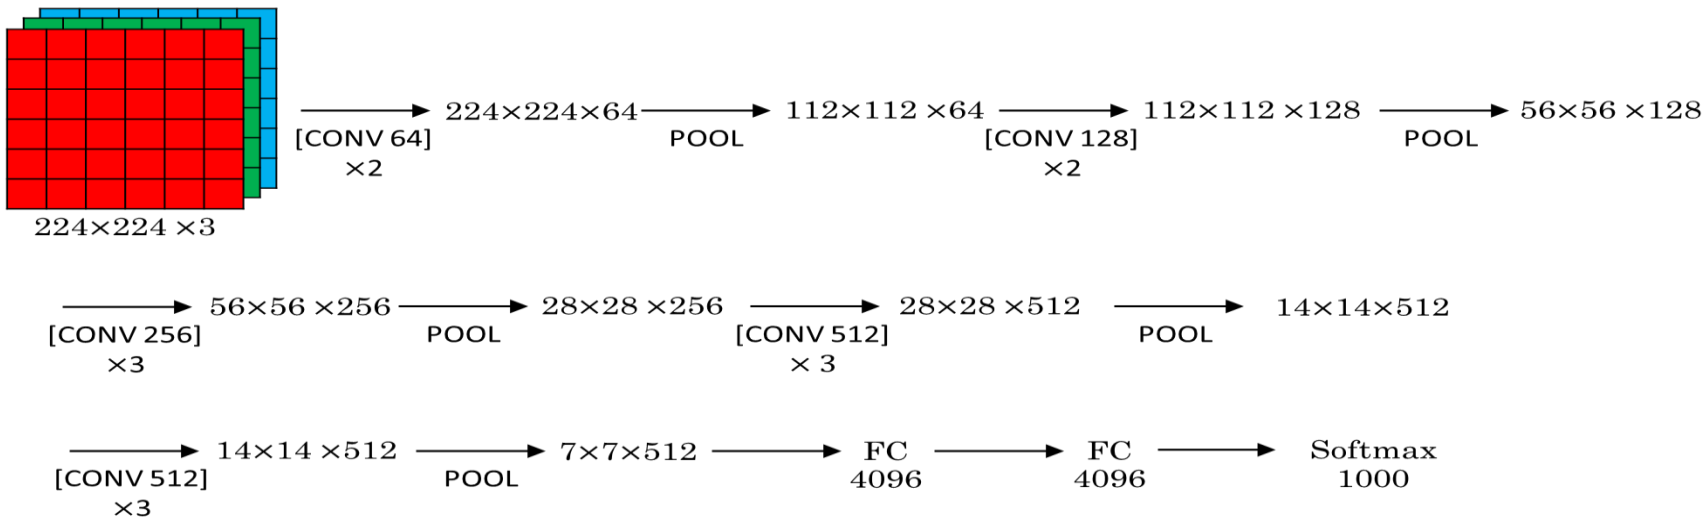
- VGG-16 model, created by the Visual Geometry Group at the University of Oxford.
- The VGG-16 model is a 16-layer (convolution and fully connected) network built on the ImageNet database, which is built for the purpose of image recognition and classification.

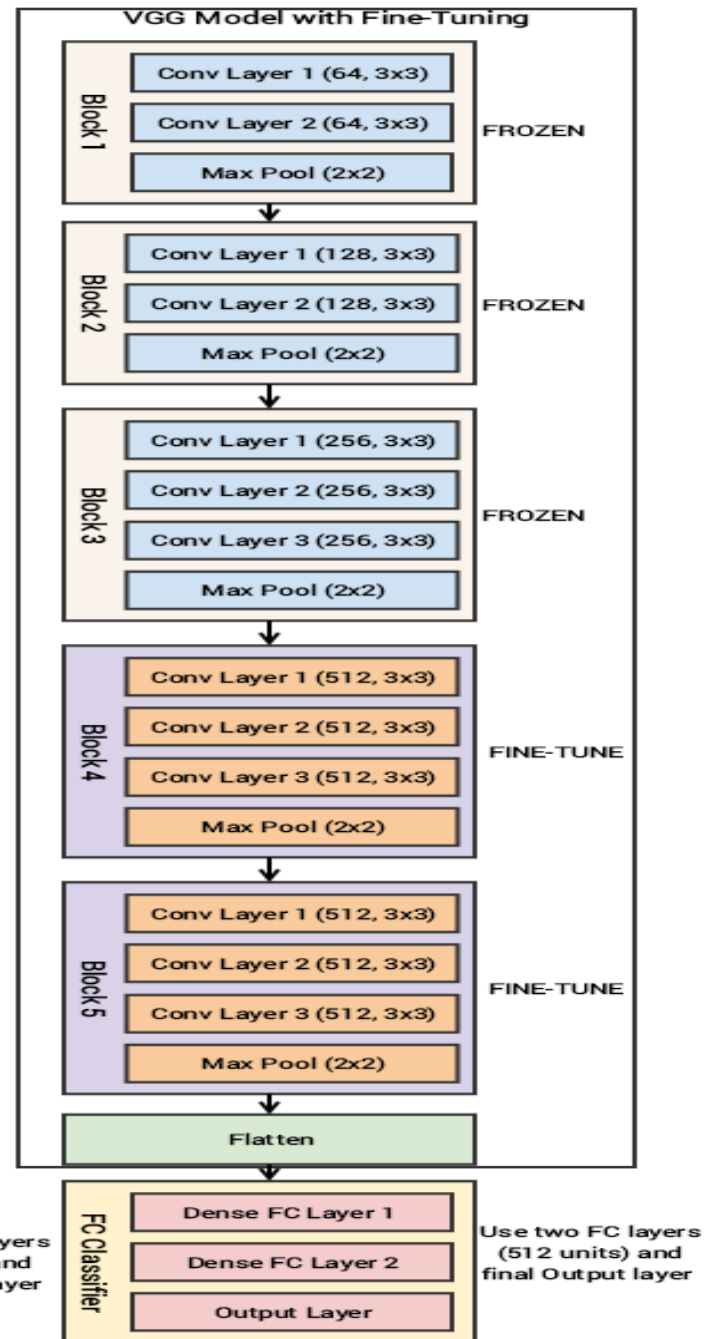
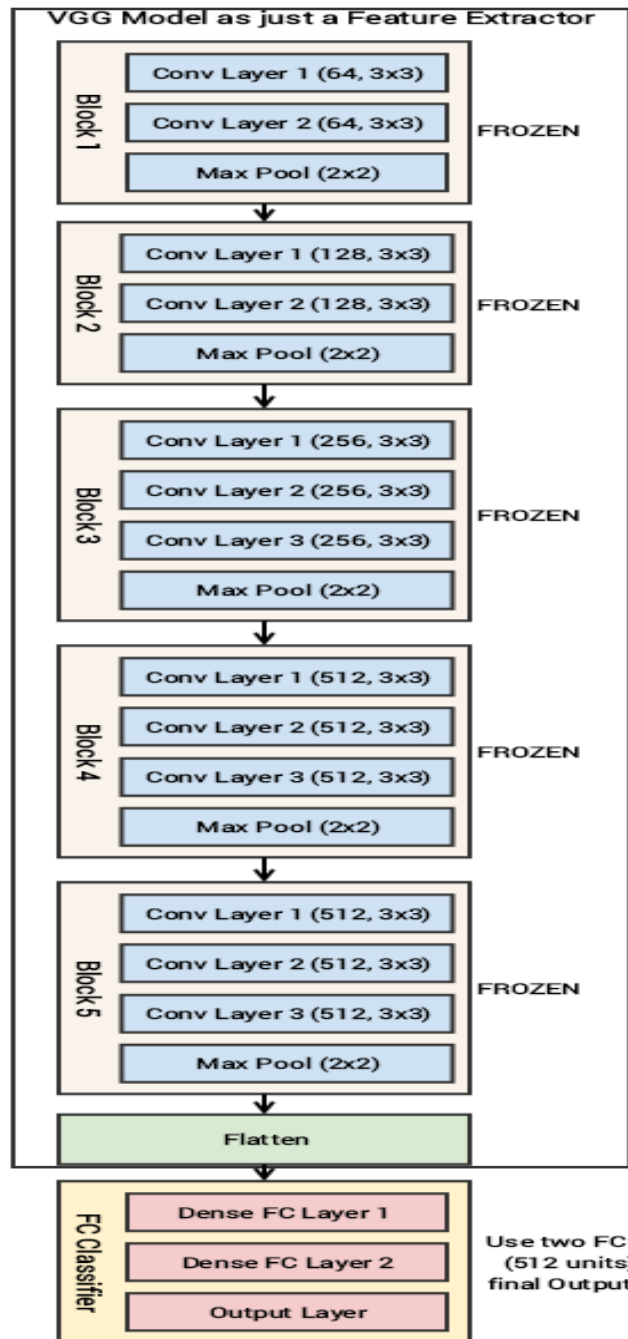



## VGG - 16

CONV =  $3 \times 3$  filter,  $s = 1$ , same

MAX-POOL =  $2 \times 2$ ,  $s = 2$





# Problems with Very Deep Networks

## •Initial Belief:


- Increasing the number of layers in a neural network would consistently improve accuracy.

## •Practical Issues:

- No generalization
- Vanishing gradient
- Exploding gradients
- Hampered training and limited performance

Vanishing Gradient

Exploding  
Gradients



# Vanishing/Exploding Gradients Problems:

In deep neural networks, during backpropagation, gradients can become extremely small. This happens as they are propagated back through the network.


## Vanishing Gradient Problem:

- During backpropagation, gradients can become extremely small.
- Lower layers' weights remain unchanged, hindering convergence..

## Exploding Gradients Problem:

- During backpropagation, Gradients grow larger leading to unstable training.
- Results in excessively large weight updates.
- Often observed in recurrent neural networks.

It makes learning difficult, particularly in deeper layers.



Comparison of 20-layer vs 56-layer architecture

# Solution

# Residual Learning

## Skip Connections:

- A shortcut connection in a neural network that bypasses some layers and directly feeds the output of an earlier layer to a later layer.
- Helps in bypassing any layer that may degrade performance, effectively regularizing the network.

## Residual Blocks:


- Introduce skip connections to address vanishing/exploding gradient issues.
- Form the basic building blocks of ResNet.

## Residual Mapping:

- Instead of learning the underlying mapping  $H(x)$ , the network learns the residual  $F(x)$  where  $F(x) := H(x) - x$
- Allows the network to fit the function  $H(x) := F(x) + x$ .

## Advantages:

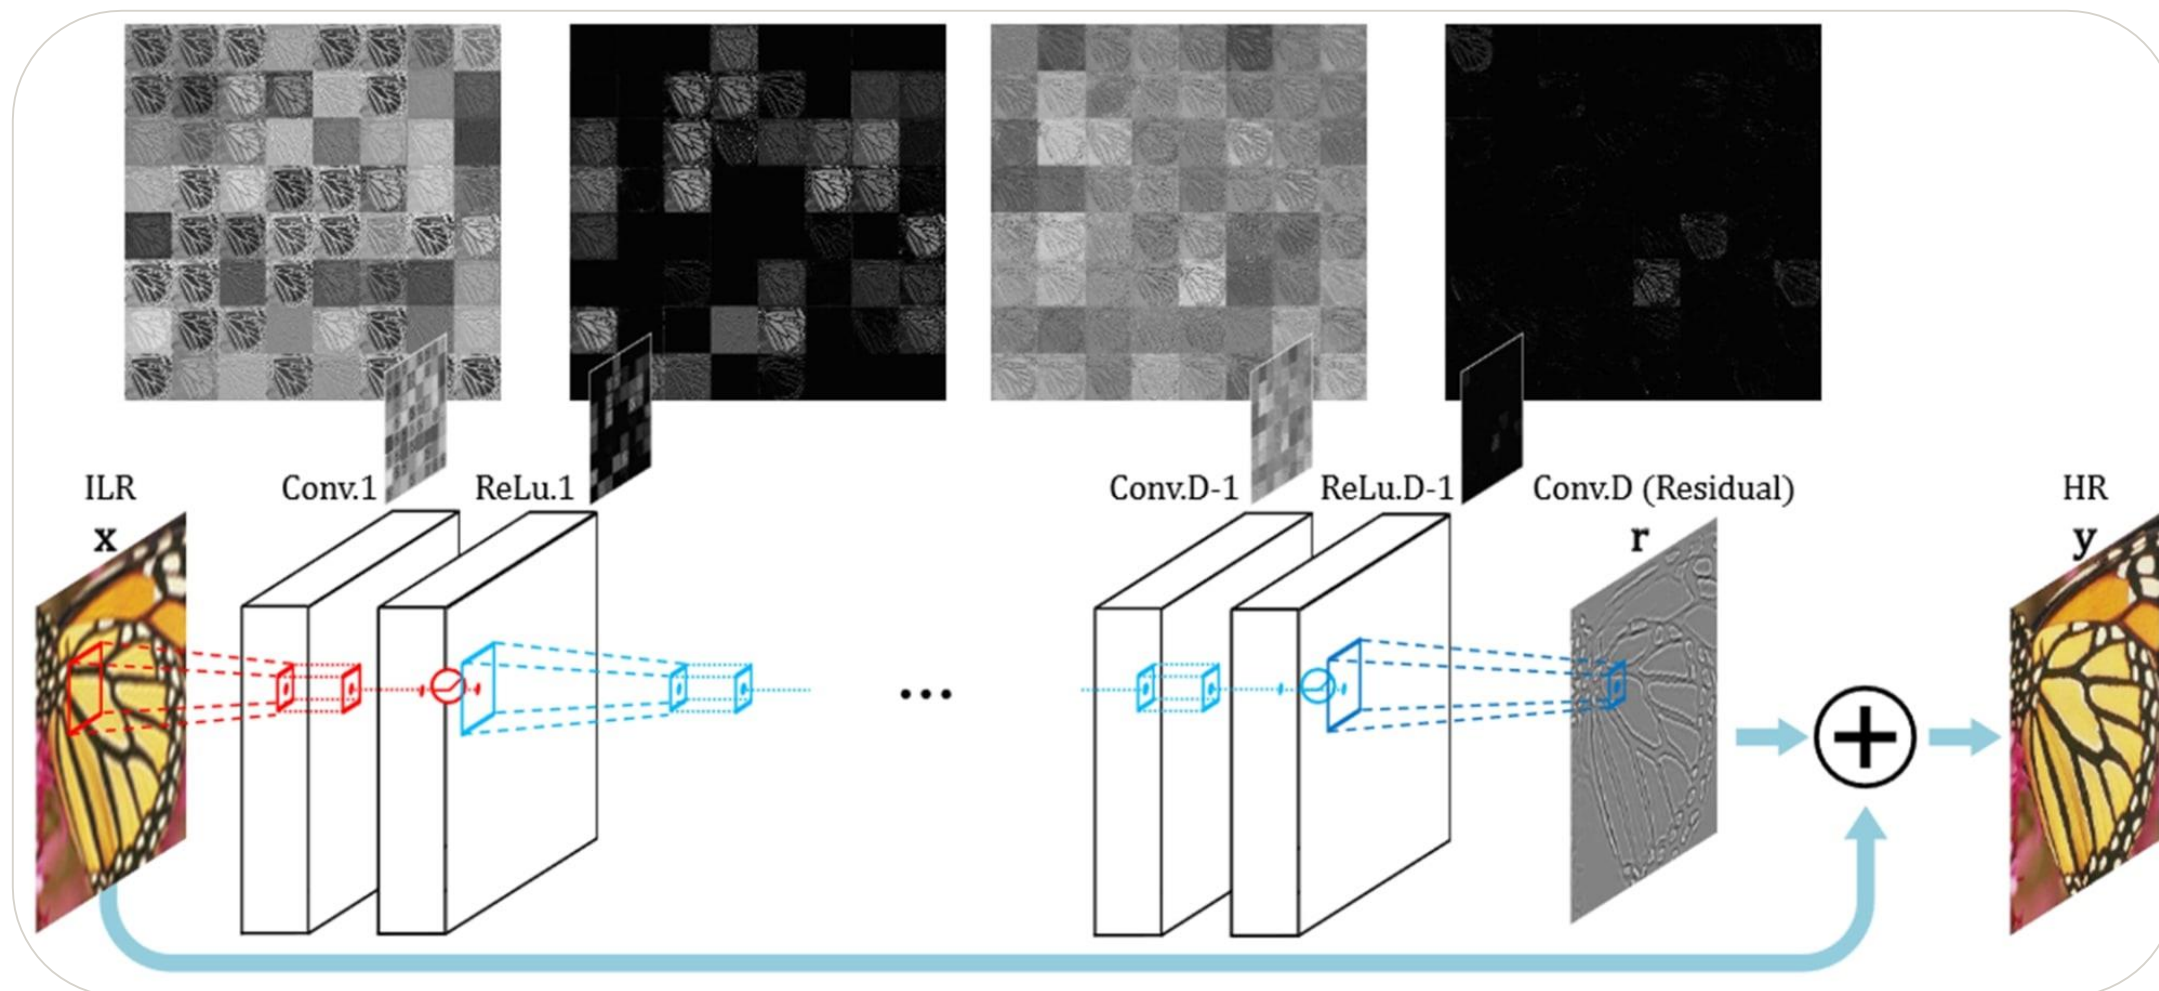
- Mitigates vanishing/exploding gradient problems.
- Enables training of very deep neural networks.
- Improves performance and generalization.



A Residual Block in a deep Residual Network. Here the Residual Connection skips two layers



# Residual Learning



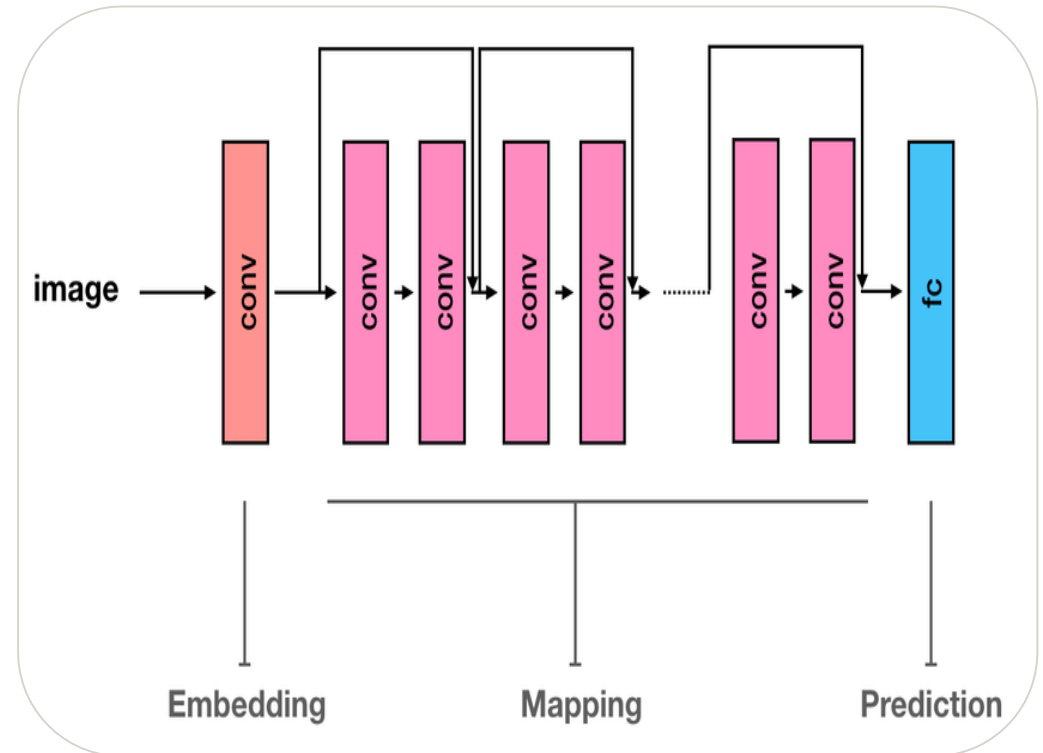


# ResNet

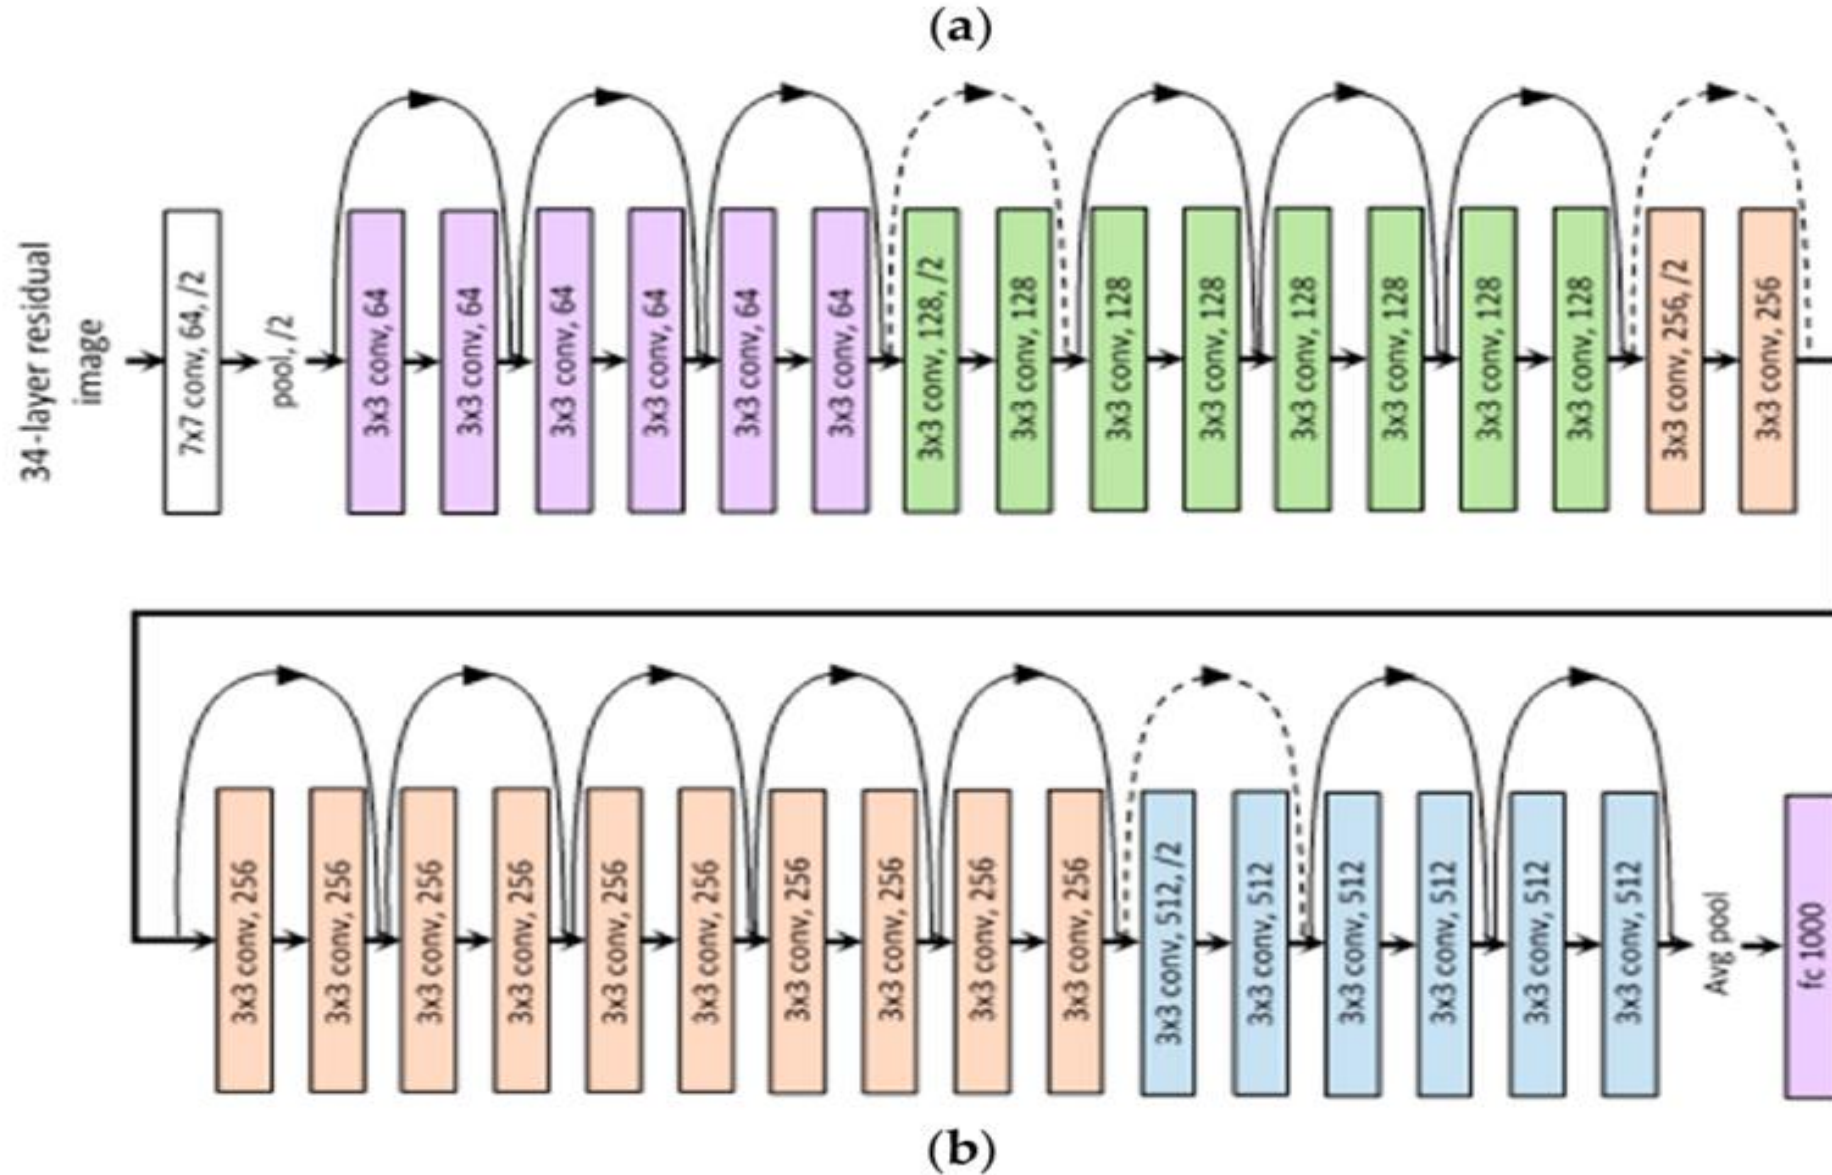
- ResNet was proposed in 2015 by researchers at Microsoft Research introduced a new architecture called Residual Network.
- A Residual Neural Network (ResNet) is an Artificial Neural Network (ANN) of a kind that stacks residual blocks on top of each other to form a network.
- ResNet first introduced the concept of skip connection.
- Winner of the ImageNet Challenge in 2015 with an error rate of 3.57%.
- ResNet50 is a variant of ResNet model which has 48 Convolution layers along with 1 MaxPool and 1 Average Pool layer

## Variants of ResNet architecture

- Resnet-18, Resnet-34, Resnet-50, Resnet-101, Resnet- 152. The number after all the model is the number of layers in the model.



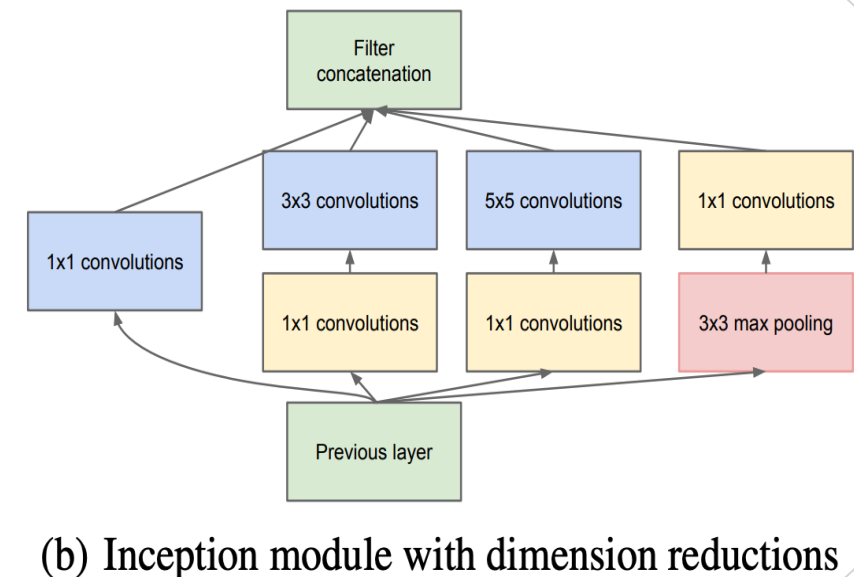
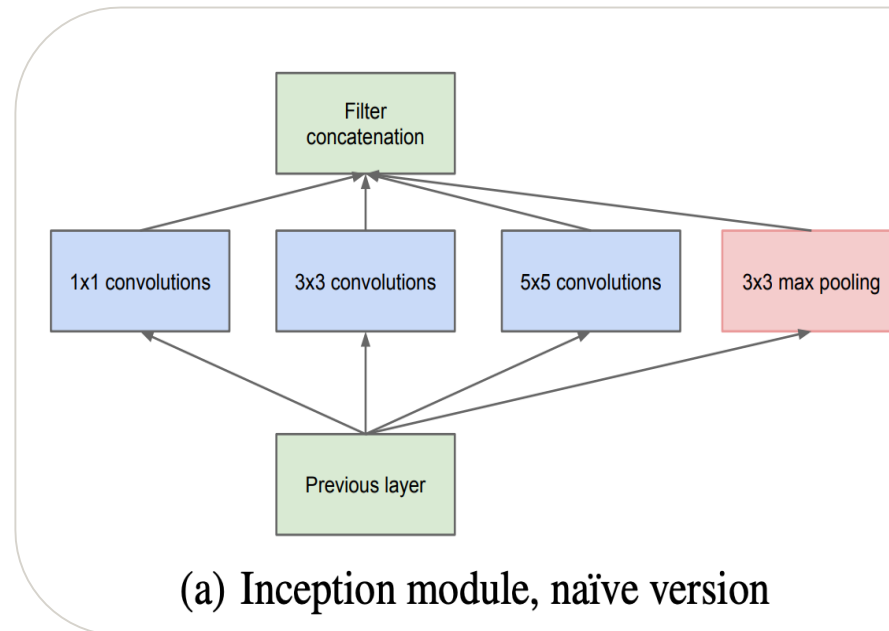

# ResNet-34 Layered architecture



# Inception

## Inception Module:

- Utilizes multiple convolutional filters (1x1, 3x3, 5x5) and pooling operations within the same module.
- While some networks like VGG16 focus only on 3x3 or LeNet5 on 5x5, Inception makes sure to grab all kinds of features.
- By using various filter sizes, Inception can pick up both small and big details in the data.
- Captures information at different scale.

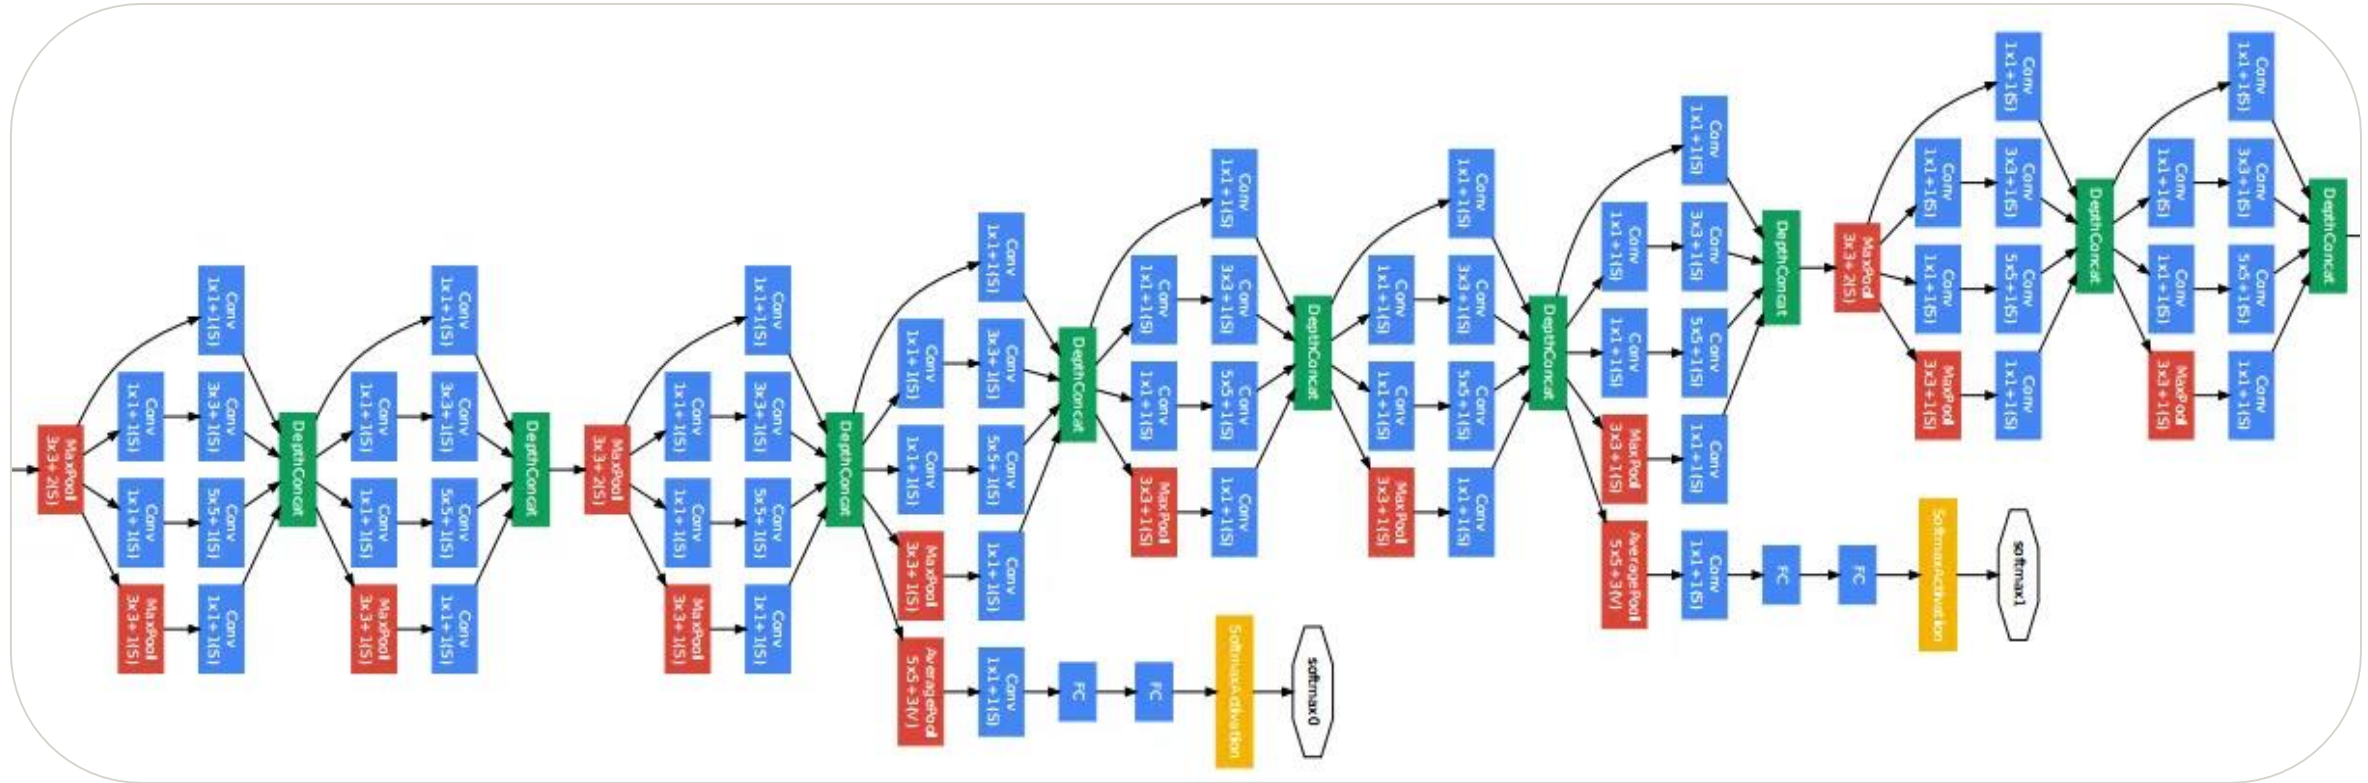


## Inception Pre-trained Models:

- Inception-v1 (GoogLeNet), Inception-v2, Inception-v3, Inception-v4.


### Notable Achievements:

- Inception-v1 (GoogLeNet) won the 2014 ImageNet Challenge with a top-5 error rate of 6.67%.



GoogLeNet, 2014

# The Problem of Computational Cost



$28 \times 28 \times \underline{192}$

→  
CONV  
 $\underline{5 \times 5},$   
same,  
 $\underline{32}$

32 filters.



$\underline{28 \times 28 \times 32}$

filters one  $\underline{5 \times 5 \times 192}$

$$\underline{28 \times 28 \times 32} \times \underline{5 \times 5 \times 192} = \underline{120M.}$$

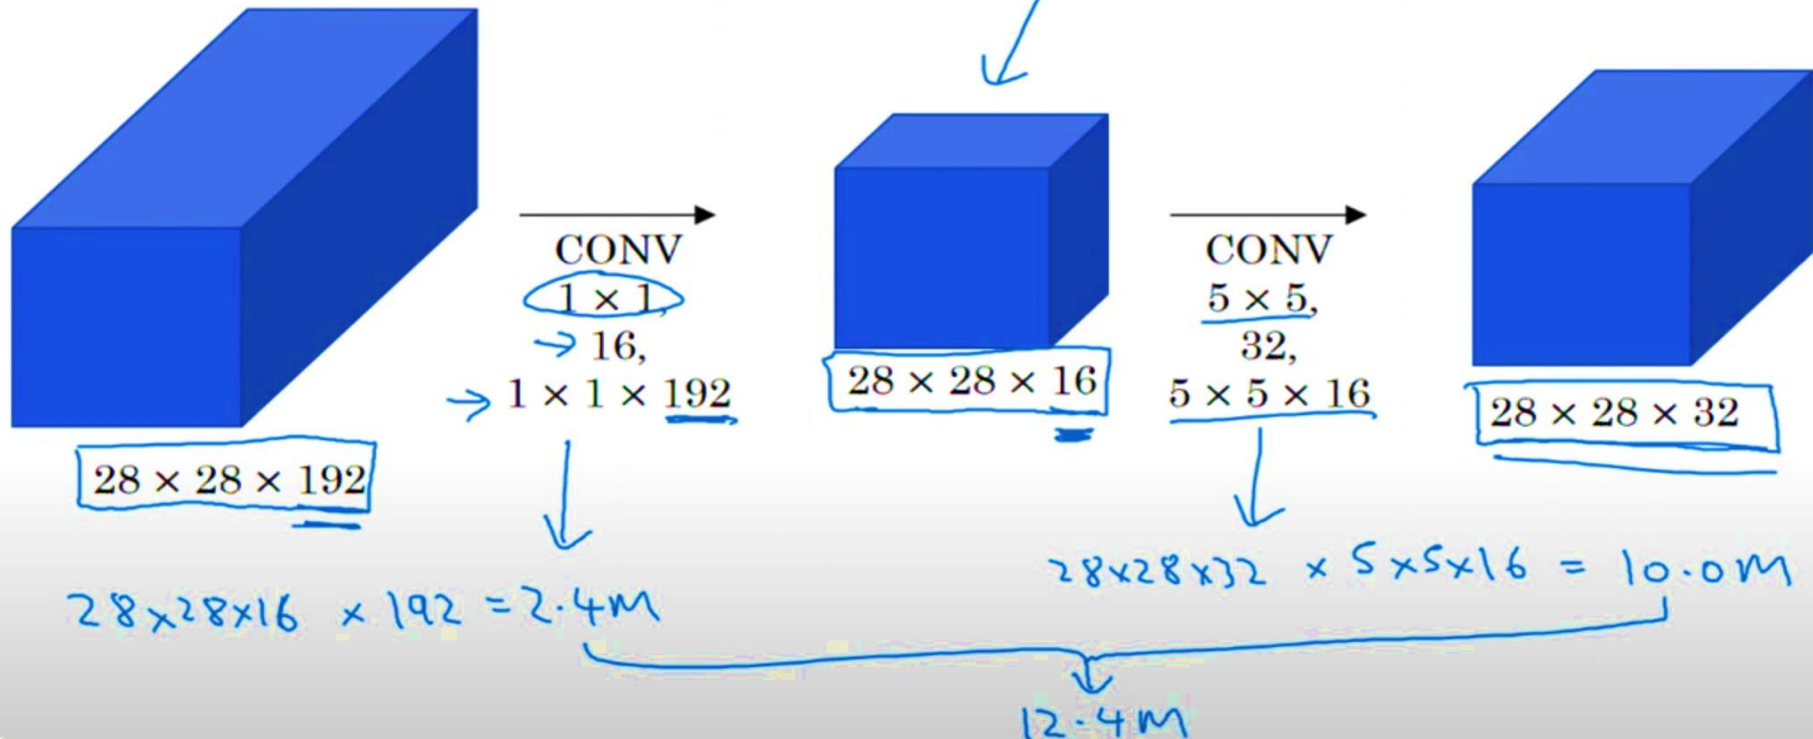


# Solution

Less Parameters means Less Computational Cost.

- Add  $1 \times 1$  Conv before  $3 \times 3$
- Add  $1 \times 1$  Conv before  $5 \times 5$
- And Add  $1 \times 1$  Conv after the  $3 \times 3$  Maxpool layer.

## Using $1 \times 1$ convolution



# 1x1 Convolution



- A 1x1 convolution applies a single 1x1 filter to each pixel in the input volume.
- It processes each pixel individually but across all channels (depth), combining the information from different channels.

## Purpose:

- Reduces the number of channels while retaining spatial dimensions.
- Enables efficient dimensionality reduction and computational cost savings.

## Applications

- **ResNet:** Used in bottleneck blocks for efficiency.
- **MobileNet:** Part of depthwise separable convolutions.
- **Inception Modules:** Reduces dimensions before expensive convolutions.



28 x 28 x 192



1 x 1 x 192 (1 kernel)



28 x 28 x 1

# EfficientNet


- Introduced in 2019 by a team of researchers at Google AI,
- The most powerful CNN architecture
- EfficientNet is built upon a concept called compound scaling.
- **Compound scaling** optimizes model depth, width, and resolution for optimal efficiency.

## Applications:

- Image classification, object detection, semantic segmentation

## EfficientNet Variants

- **EfficientNet B0-B7**: A family of EfficientNet models with varying complexities.
- **B0**: Most lightweight (5.3 million parameters).
- **B7**: Most complex (6.1 billion parameters).





# MobileNet

Developed by Google researchers.

## Purpose:

- Designed for mobile and embedded vision applications.
- Focuses on efficient, lightweight models suitable for devices with limited computational resources.

## Key Features:

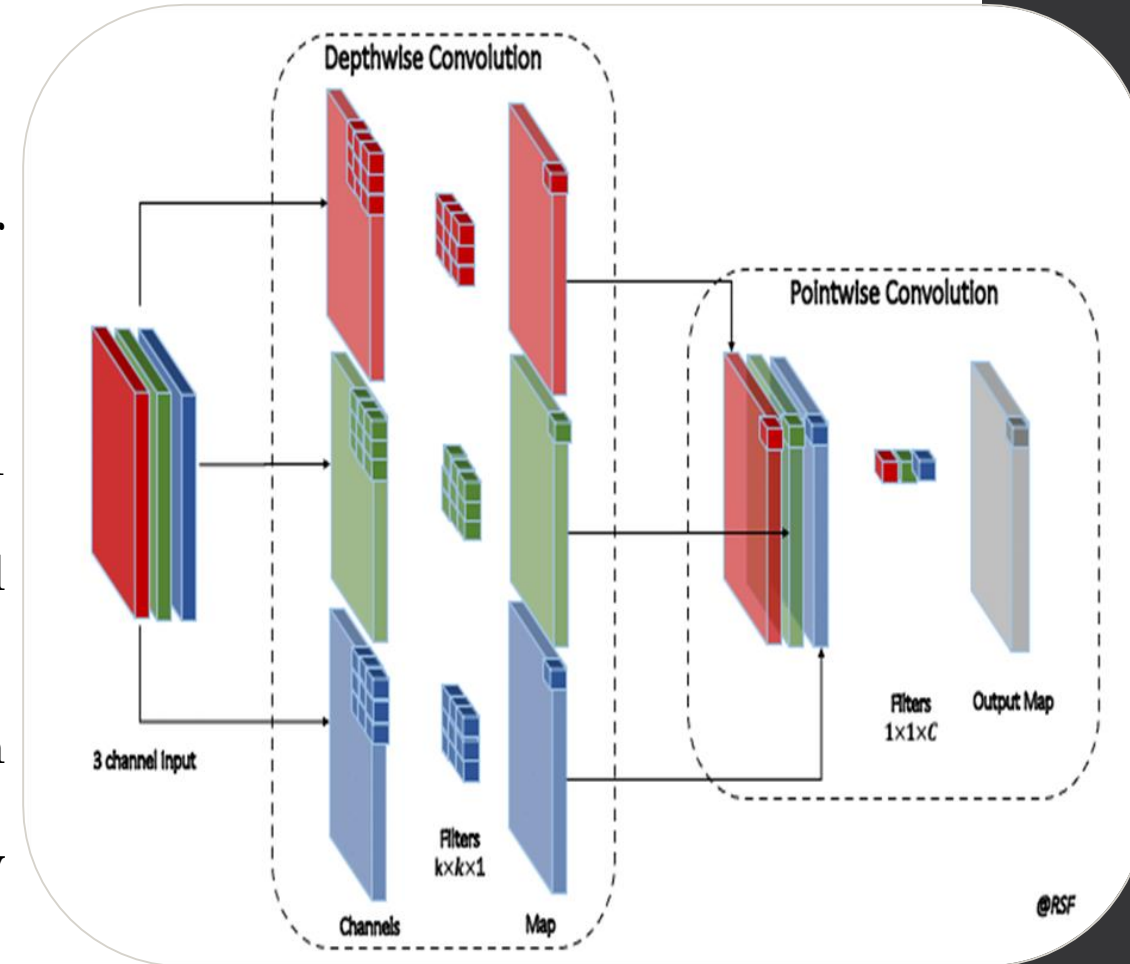
- Reduces computational cost and model size.
- Fewer parameters compared to traditional convolutional networks.
- Maintains competitive accuracy with optimized speed and efficiency.

## Applications:

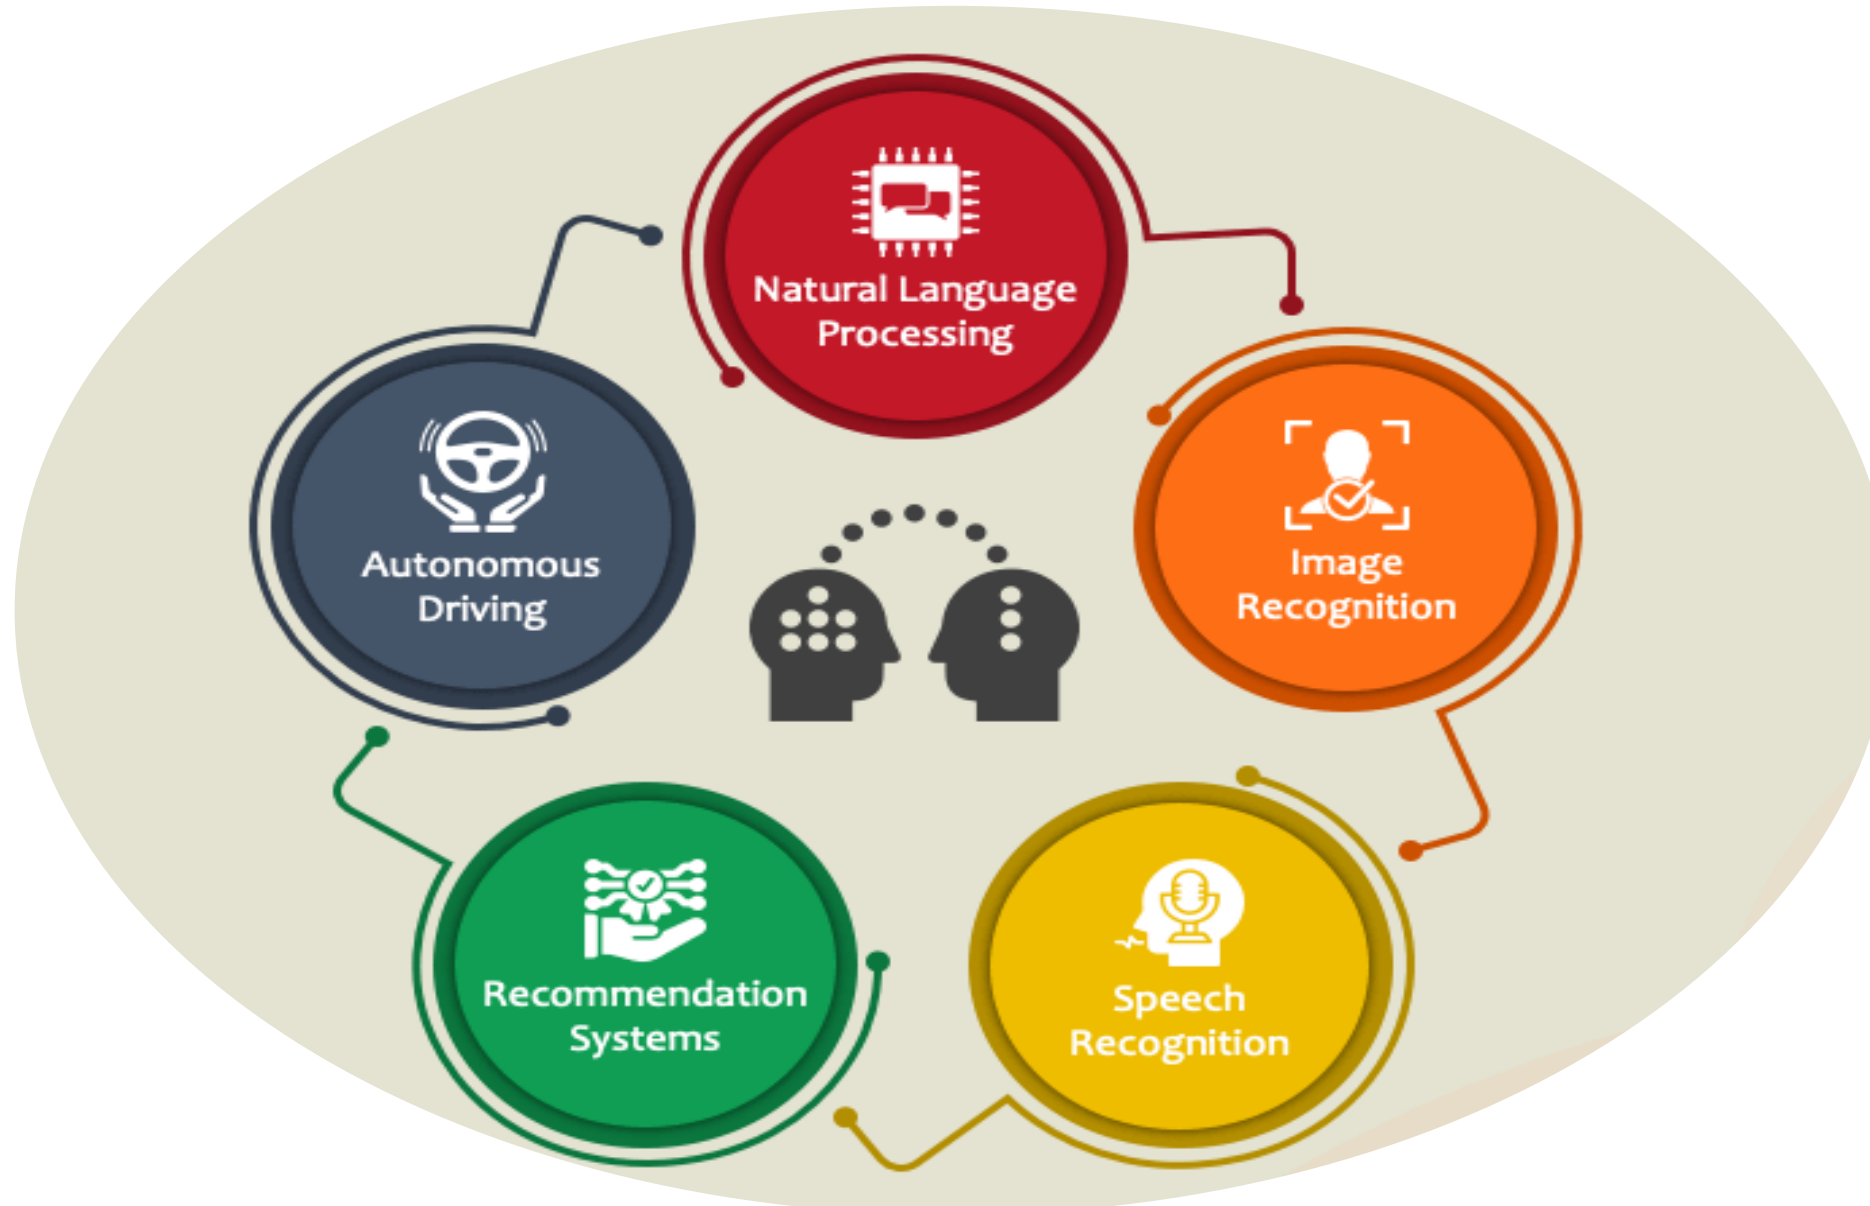
- Real-time object detection and image classification on mobile devices.
- Deployment in IoT devices and augmented reality applications.

## Pre-trained Models:

- MobileNetV1, MobileNetV2, MobileNetV3.
- Pre-trained on ImageNet, available for transfer learning.



# Applications of transfer Learning



# Applications of transfer Learning

## Image Classification

A core application of transfer learning in computer vision.

### Pre-trained Models

- Leverage powerful models like ResNet, VGG, and Inception.
- Trained on massive datasets like ImageNet.
- Fine-tune models for specific domains.

### Applications:

- For examples, Identifying species in wildlife photography or diagnosing medical conditions from imaging data.

## Object Detection:


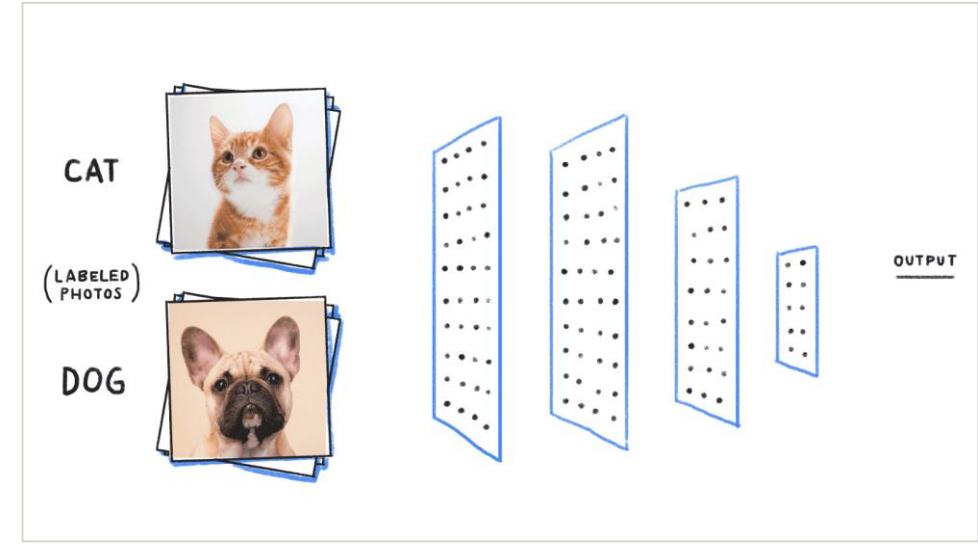
Detect and localize objects in images or videos.

### Pre-trained Models:

- Utilize pre-trained models like YOLO, Faster R-CNN, SSD for feature extraction.
- Add layers for bounding box prediction and class identification.

### Applications:

- Pedestrian detection for self-driving cars.





# Applications of transfer Learning

## Image Segmentation

Segmenting images into distinct regions corresponding to objects or parts of objects.

### Pre-trained Models:

- U-Net, DeepLab, FCN.

### Applications with Transfer Learning

- **Medical Imaging:** Identify tumors or other abnormalities.
- **Autonomous Vehicles:** Differentiate between roads, sidewalks, and vehicles.

### Face Recognition:


Identify and verify faces in images or videos.

### Pre-trained Models:

- Utilize pre-trained models like FaceNet, VGGFace for feature extraction.
- Add layers for face identification and verification.

### Applications:

- Security systems for access control.
- Social media tagging.
- User authentication for devices and apps.



## References:

- [https://www.researchgate.net/figure/Change-in-the-number-of-model-parameters-and-training-time-with-increased-number-of\\_fig3\\_335865710](https://www.researchgate.net/figure/Change-in-the-number-of-model-parameters-and-training-time-with-increased-number-of_fig3_335865710)
- <https://blog.devgenius.io/all-you-need-to-know-about-transfer-learning-4a7e3cbaf1dd>
- <https://medium.com/@nutanbhogendrasharma/transfer-learning-using-feature-extraction-in-deep-learning-afd97380c96c>
- [https://dev.mrdbourke.com/tensorflow-deep-learning/04\\_transfer\\_learning\\_in\\_tensorflow\\_part\\_1\\_feature\\_extraction/](https://dev.mrdbourke.com/tensorflow-deep-learning/04_transfer_learning_in_tensorflow_part_1_feature_extraction/)
- <https://encord.com/glossary/pre-trained-model-definition/>
- <https://en.wikipedia.org/wiki/ImageNet>
- <https://blog.mturk.com/tutorial-how-to-label-thousands-of-images-using-the-crowd-bea164ccbefc>
- <https://medium.com/@siddheshb008/alexnet-architecture-explained-b6240c528bd5>
- 62.3 million learnable parameters.
- <https://www.philschmid.de/getting-started-with-cnn-by-calculating-lenet-layer-manually>
- <https://medium.com/@mygreatlearning/everything-you-need-to-know-about-vgg16-7315defb5918>
- <https://www.geeksforgeeks.org/residual-networks-resnet-deep-learning/>
- <https://viso.ai/deep-learning/resnet-residual-neural-network/>
- [https://www.researchgate.net/figure/A-schematic-view-of-ResNet-architecture-15-decomposed-into-three-blocks-embedding\\_fig1\\_333475917](https://www.researchgate.net/figure/A-schematic-view-of-ResNet-architecture-15-decomposed-into-three-blocks-embedding_fig1_333475917)

## References:

- <https://medium.com/@siddheshb008/resnet-architecture-explained-47309ea9283d>
- <https://velog.io/@snoop2head/Going-Deeper-with-Convolutions-GoogleNet-Inception>
- <https://towardsdatascience.com/an-intuitive-guide-to-deep-network-architectures-65fdc477db41>
- <https://www.youtube.com/watch?app=desktop&v=pf-HUcqdCr4>
- [https://www.baeldung.com/wp-content/uploads/sites/4/2020/06/3D\\_1D\\_cropped.gif](https://www.baeldung.com/wp-content/uploads/sites/4/2020/06/3D_1D_cropped.gif)
- <https://arxiv.org/abs/1905.11946>
- <https://deeplobe.ai/image-segmentation-the-most-interesting-applications/>
- <https://medium.com/analytics-vidhya/how-facial-recognition-systems-work-edcbb227e614>

Thank You